

Reviewing the prompt

This is where most send-backs come from, because the prompt can't be edited. Read it as a real person would, then solve it yourself against the environment files before judging anything else.

Always solve the prompt first

You can't tell whether a model failed - or whether a justification is right - until you know the correct answer from the files. Work the prompt end to end before you read the responses.

What a good prompt has

Work through each item.

Checklist

0 of 4 complete

☐

Clearly uses the assigned persona

The persona's occupation, income, family, and spending show through in the question.

☐

Sounds like a real person

Natural, first-person phrasing - the kind of question someone actually types, not a templated paragraph.

☐

No LLM-generated text

Strictly not allowed. Machine-written prompts get sent back.

☐

A real ask that still breaks models

A genuine personal-finance question that is meaningfully difficult and produces real model failures.

What counts as a meaningful model failure

A meaningful model failure is something the model got **objectively wrong** in a way that changes the answer or violates the prompt. Cosmetic differences, style preferences, and harmless variations don't count.

Hallucination

Counts - the model invented a charge, account, or figure that isn't in the files.

Doesn't count - the model used a slightly awkward synonym for a real figure.

Miscalculation

Counts - wrong total, wrong rate, or a dropped transaction.

Doesn't count - a rounding difference in the last cent that doesn't change the answer.

Files and data

Counts - the model couldn't find the statements, or fabricated data not in the workspace.

Doesn't count - the model opened the files in a different order than you would.

Workflow and send-backs

Your default is to fix and improve the task, not reject it. Most issues are minor and fixable in place. Send-backs are reserved for a small set of real failures.

The one-line rule
Fix what you can. Send back only a broken prompt, a task with no real model failure, or a rubric too poor to salvage. Everything else is a minor edit you make yourself.

When to fix vs. Send back

Situation	Action	Avoid
Rubric wording is sloppy	Fix it - tighten the criterion, correct the weight, move on.	Sending the whole task back over an editable rubric line.
A justification is vague but the failure is real	Fix it - sharpen the justification with the exact numbers.	Rejecting a task that has a genuine model failure.
The prompt doesn't resolve to a correct answer	Send it back - prompts can't be edited.	Trying to patch the prompt yourself, or approving it anyway.
No model actually failed	Send it back - there's nothing to deliver.	Manufacturing a failure to clear the difficulty bar.

Hard rule - prompt problems are always a send-back
Prompts cannot be edited by reviewers. If the prompt is broken - wrong, ambiguous, off-persona, or it doesn't resolve against the files - the only path is to send it back to the fellow with feedback. You can never fix a prompt in place.

Your review order

- 1

The prompt
Persona, natural phrasing, no LLM-written text, and - most important - it resolves to a single correct answer from the environment files.
- 2

The model responses and justifications
Confirm the failures are meaningful, not minor, and that each justification's numbers are actually correct against the files.
- 3

The rubric
Coverage, no double jeopardy, self-contained, verifiable and objective, atomic.

✓ Good

"Response states that a down payment under 20% triggers PMI."

⚠ Bad

"Response considered the user's PMI exposure."

"Considered" is mind-reading. You can only grade what's written in the response.

Common error indicators - fix criteria that use these

Process words (describe how, not the result): calculates, computes, derives, determines.

Vibe / non-objective words (two graders would disagree): good, clear, helpful, appropriate, reasonable, sufficient, nice, well-structured, properly, adequate, suitable, thorough, accurate, correctly, comprehensive, robust, sound.

Mind-reading words (can't see inside the model): considered, understood, took into account, thought about, analyzed, evaluated, recognized, realized, knew, accounted for.

The fix is almost always the same: rewrite the criterion to state the final value or outcome plainly
- "Response states X."

E. Atomic

Each criterion checks **1 thing**, so it can't be half-right.

✓ Good

"(1) Response states the monthly gap is \$1,654.40. (2) Response uses the requested four section headings: What You Actually Have, Monthly Target, Is It Realistic, What To Watch."

⚠ Bad

"Response states the monthly gap is \$1,654.40 and uses the requested four section headings: What You Actually Have, Monthly Target, Is It Realistic, What To Watch."

If two parts could independently pass or fail, split them. Apply with judgment - rubrics are short (20–35 criteria), so don't fragment past the point of usefulness.

Example - Denny (joint statements, provable vs. suspicious duplicates)

"Response correctly classifies duplicates" is not atomic - it bundles five separate calls. Split into: (1) Netflix (\$22.99) classified as provable duplicate, (2) Orangetheory (\$159.00) classified as provable duplicate, (3) Verizon (\$165.00) classified as suspicious-not-proven, (4) Progressive (\$185.00) classified as suspicious-not-proven, (5) Provable-duplicate total stated as \$182.00/month.

C. Self-contained

Every criterion must be gradable from its own text alone - no references to anything outside it, especially no file names.

✔ Good

"Response states the Bank of America checking ending balance is \$29,450.67."

⚠ Bad

"Response states the credit card statement value from march_statement.pdf."

A grader with no access to the prompt or the files must be able to check the criterion. Name the final value, not the source.

Common self-contained failures from real tasks

"Response uses the figure from `priya_vasanthan__bofa_checking__202604` ." → name the figure: \$29,450.67.

"Response correctly identifies the duplicates listed in the prompt." → name them: Netflix (\$22.99) and Orangetheory (\$159.00).

"Response matches the SEP contribution limit referenced above." → name it: \$13,535.09 for 2026.

D. Verifiable and objective

A criterion must check an outcome visible in the output - not how the model got there, and not its internal reasoning.

✔ Good

"Response states the monthly card-balance gap is \$1,654.40."

⚠ Bad

"Response calculates the monthly card-balance gap."

"Calculates" describes the model's process. The grader can't see process - only the final value.

✔ Good

"Response states the average monthly credit-card spend is \$5,319.50."

⚠ Bad

"Response gives a clear and helpful budget breakdown."

The good version names the exact values a grader can verify against the response. "Clear" and "helpful" are vibe words - two graders would disagree on whether the bar was met.

Reviewing the rubric

Rubrics run 20–35 criteria, each with a weight from –100 to +100 and a one-line rationale. Check the rubric on 5 things - every example below is a real criterion shape pulled from the golden tasks.

A. Coverage

The rubric must capture the **most important things the prompt asks for**. If a reasonable, correct answer would score well and a bad answer would score poorly, coverage is good. Read the prompt's real asks and confirm each maps to a criterion.

Example - Priya (homebuyer feasibility, 36 months)

The prompt's headline ask is "can I hit a \$165,000 down payment in three years?" Coverage means the rubric has criteria for: (1) the real spendable cash after the \$22,795.95 unpaid tax, (2) the \$2,499.84/month savings target, (3) the explicit 401(k) exclusion, and (4) the four-section deliverable format. Miss any one of those and the rubric does not cover the prompt.

✓ Good

"Response states the monthly savings target to reach \$165,000.00 in 36 months is \$2,499.84."

⚠ Bad

"Response calculates a monthly savings target."

The "Good" version locks in the prompt's headline number. The "Bad" version doesn't cover it - a model that returns \$1,866.62 would still pass.

B. No double jeopardy

Don't penalize the same thing twice

Watch for two criteria checking the same thing, and for a positive/negative pair that mirror each other. If "Response identifies the avalanche method" is +80, there must not also be a "Response fails to identify the avalanche method" at –80. The zero points from a missed positive is already the penalty.

✓ Good

"+60 - Response states the real spendable cash is \$75,005.92 (after the \$22,795.95 tax liability)."

⚠ Bad

"+60 - Response identifies the unpaid tax liability of \$22,795.95. AND –60 - Response fails to subtract the \$22,795.95 tax liability from spendable cash."

Both "Bad" criteria score the same underlying finding. Keep one, and let the zero do the work on a miss.

The 2 most common prompt-section problems

1.The prompt doesn't resolve to a correct answer from the files



2.The justification is wrong, so the failure is false



"Sometimes the model actually answered correctly, and the justification arrives at different values that aren't the real answer. That isn't a model failure - it's a mistake in the justification. Verify every number in the justification against the files before you trust that a model failed."

Answer

Counts - the model reached no valid answer, or no answer at all.

Doesn't count - the model reached the right answer but formatted it differently than you'd prefer.

Instructions

Counts - the model explicitly violated a constraint the prompt stated.

Doesn't count - the model did something the prompt never asked it not to do.

The 2 most common prompt-section problems

1. The prompt doesn't resolve to a correct answer from the files

"Work the prompt yourself against the documents. If you can't reach a single defensible answer - because the files contradict each other, a needed number is missing, or the question is ambiguous - the prompt is broken. Send it back. A failure that comes from an ambiguous prompt is the author's fault, not the model's."

2. The justification is wrong, so the failure is false
